

# Bayesian multimodeling: diffusion

2024

# What is diffusion?

A process  $X_t$  is a diffusion if it has these properties:

- The time variable  $t$  is continuous;
- It is a Markov process;
- $X_t$  is a continuous function of  $t$ .


A diffusion process  $X_t$  can be expressed in Ito form:

$$dX_t = a(X_t)dt + b(X_t)dW_t,$$

where  $a$  is a drift,  $b$  is a noise coefficient,  $W_t$  is a Wiener process.

# Diffusion forward process

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I}).$$




$$\alpha_t = \prod \beta_t, \quad q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\sqrt{\alpha_t}, (1 - \alpha_t)\mathbf{I}),$$

$\beta_t$  is set for  $\alpha_t \rightarrow 0$ , and  $q(\mathbf{x}_t|\mathbf{x}_0) \approx \mathcal{N}(\mathbf{0}, \mathbf{I})$ .

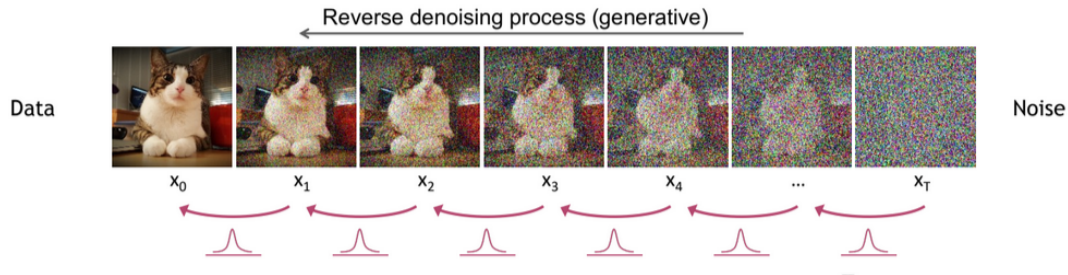
# Diffusion forward process

## Diffused Data Distributions



# Diffusion backward process

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\boldsymbol{\mu}(\mathbf{x}_t), \sigma^2 \mathbf{I}).$$



# Diffusion as variational model

ELBO:

$$\begin{aligned} \mathbb{E}_q \log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} &= \\ &= \mathbb{E}_q \log p_\theta(\mathbf{x}_0|\mathbf{x}) - \sum_t KL(q(\mathbf{x}_{t-1}|\mathbf{x}_t)|p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)) + KL(q(\mathbf{x}_T|\mathbf{x}_0)|p_\theta(\mathbf{x}_T)). \end{aligned}$$

Sum of KL can be computed by explicit formula:

$$KL(q(\mathbf{x}_{t-1}|\mathbf{x}_t)|p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)) = \text{Const} + \mathbb{E}_q \frac{1}{2\sigma^2} \left\| \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1-\alpha_t}} \boldsymbol{\varepsilon} \right) - \boldsymbol{\mu}(\mathbf{x}_t) \right\|^2.$$

# Diffusion model training

---

## Algorithm 1 Training

---

```
1: repeat  
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$   
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$   
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
5:   Take gradient descent step on  
        $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$   
6: until converged
```

---

---

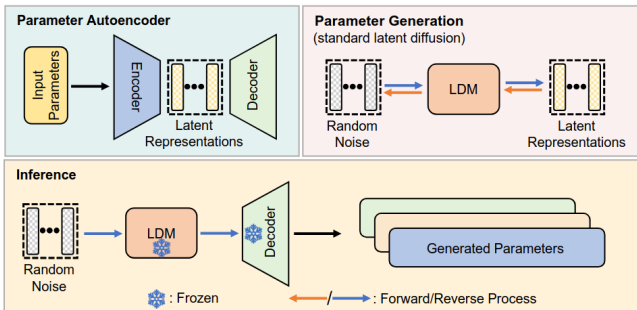

## Algorithm 2 Sampling

---

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
2: for  $t = T, \dots, 1$  do  
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$   
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$   
5: end for  
6: return  $\mathbf{x}_0$ 
```

---

# Model parameters generation using diffusion models





# Hypernetworks


## Definition

Given a set  $\Lambda$ .

Hypernetwork is a parametric mapping from  $\Lambda$  to set  $\mathbb{R}^n$  of the model  $\mathbf{f}$  parameters:

$$G : \Lambda \times \mathbb{R}^u \rightarrow \mathbb{R}^n,$$

where  $\mathbb{R}^u$  is a set of hypernetwork parameters.



# Diffusion and score matching

Score matching idea:

$$p_{\theta}(\mathbf{x}) = \frac{1}{Z(\theta)} \exp(-E(\mathbf{x}; \theta)).$$

To approximate unknown true distribution  $p(\mathbf{x})$  we try to match scores of the distribution:

$$\begin{aligned} & \mathbb{E}_{\mathbf{x} \in p(\mathbf{x})} \left\| \frac{\partial \log p_{\theta}(\mathbf{x})}{\partial \mathbf{x}} - \frac{\partial \log p(\mathbf{x})}{\partial \mathbf{x}} \right\|^2 \rightarrow \min. \\ & \approx \mathbb{E}_{\hat{\mathbf{x}}, \mathbf{x} \in q(\hat{\mathbf{x}}, \mathbf{x})} \left\| \frac{\partial \log p_{\theta}(\mathbf{x})}{\partial \mathbf{x}} \Big|_{\hat{\mathbf{x}}=\mathbf{x}} - \frac{\partial \log p(\hat{\mathbf{x}}|\mathbf{x})}{\partial \hat{\mathbf{x}}} \right\|^2 \rightarrow \min, \end{aligned}$$

where  $q(\hat{\mathbf{x}}) = q(\hat{\mathbf{x}}|\mathbf{x})p(\mathbf{x})$ .

# Autoencoder: generative model?

(Alain, Bengio 2012): consider regularized autoencoder:


$$||\mathbf{f}(\mathbf{x}, \sigma) - \mathbf{x}||^2,$$

where  $\sigma$  is a noise level.

Then

$$\frac{\partial \log p(\mathbf{x})}{\partial \mathbf{x}} = \frac{||\mathbf{f}(\mathbf{x}, \sigma) - \mathbf{x}||^2}{\sigma^2} + o(1) \text{ when } \sigma \rightarrow 0.$$

Vector field induced by reconstruction error



# Reference

- Lectures on Diffusion processes:  
<https://math.nyu.edu/~goodman/teaching/StochCalc2018/notes/Lesson2.pdf>
- Tutorial on diffusion models: <https://cvpr2023-tutorial-diffusion-models.github.io/>
- Wang, Kai, et al. "Neural network diffusion." arXiv preprint arXiv:2402.13144 (2024).
- Vincent, Pascal. "A connection between score matching and denoising autoencoders." Neural computation 23.7 (2011): 1661-1674.
- Alain G., Bengio Y. What regularized auto-encoders learn from the data-generating distribution // The Journal of Machine Learning Research. – 2014. – T. 15. – №. 1. – C. 3563-3593.

## Organizational issues

- Technical meeting: next week, regular time
- Format: each team shows the basic code. You should have an idea how this will be transformed into demo.
- The draft version of blog-post and documentation will be checked in offline-mode.
- All the presented materials must be stored at the github
  - ▶ for blog-post, you can put read-only link for the overleaf if you write the post here.
  - ▶ Time-limit: 10 minutes.

# Blog-post

Minimum to write:

- Abstract/Motivation text
- Plan with all the sections you are planning and a small explanation text for each section
- Add all the planned figures or their description (*"There will be a scatter plot with our benchmark of algorithms. The x axis corresponds to the accuracy, y axis corresponds to the model size"*).
- Think and add a figure/description of the figure for the headline.
- The text is in English.
- This is a minimum, the more you have the better it is.

# Documentation

- The documentation must be available to deploy:
  - ▶ either it must be available online
  - ▶ Or there must be instructions how to deploy it in stand-alone mode
  - ▶ **At the end of the project the documentation must be deployed online.**
- The structure must be established (have a look at your favorite libraries).
- The documentation must be able to run (at githubpages or standalone on your PC).
- For the demo show at least one auto-doc.
- Popular engines: mkdocs, sphinx.



# Unit-tests

- Cover as much functions (and instructions in them) as possible. Minimal coverage is 75%.
- Synthetic tests are wellcome: generate dataset on the fly and check that at least your functions are working on them correctly.
- Segmented code is easier to test: try to segment you code into more or less elementary functions.
- Popular engines: unittest (built-in into python) + coverage, pytest + pytest-cov.