Bayesian multimodeling: graphical models

2024

Graphical models

Conditional independence

Events X, Y are conditionally independint w.r.t. Z : $\quad X \perp Y | Z$, if

$$P(X|Y,Z) = P(X|Z).$$

Conditional dependence

Events X, Y are conditionally dependent w.r.t. $\mathfrak{S}$: $X, Y \in \mathfrak{S}$, if

$$X \not\perp Y | \mathfrak{S} \setminus \{X, Y\}.$$

Graphical models

A probabiliy model is graphical, if it can be represented as a graph, where the edges correspond to conditionally dependent events.

Non-graphical models

- MLP, decision trees, etc.
- Models with complex behaviour:

$$Y \ X1 * X2 * X3 + X1 * X2 + X1 * X3 + X2 * X3$$

Types of graphical models

- Directed models (aka Bayesian networks)
  - Easy to desing
- Undirected (Markov models)
- Factor-graphs
  - Easy to infer and optimize

Plate notation

Plate notation is an alternative visuzliation for graphical models.

Elements:

- White circles (random variables);
- Grey circels(observed variables);
- Small circles (deterministic values);
- Plates (batching).

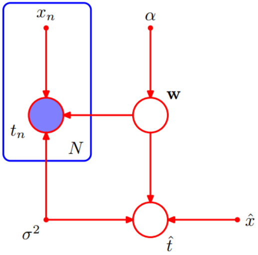

Plate notation for linear regression (Bishop)

Bayesian networks

- Models are set using directed acyclic graphs
- Joint distribution for the graph with K vertices:

$$p(v_1, \ldots, v_k) = \prod_{i=1}^{K} p(v_i | \text{parent}(v_i))$$

- Example: linear regresssion

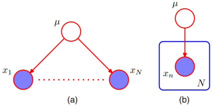

DAG and Plate notation (Bishop)

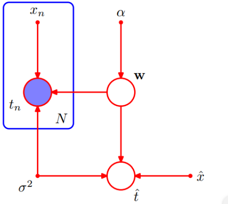

Plate notation for regression model (Bishop)

Causality graph elements

$$X \to Y \to Z - \textbf{chain}$$

Example:

- X — school budget
- Y — average student score
- Z — unviersity acceptance ratio

Properties:

1. X and Y, Y and Z are dependent:
   $\exists x, y : \mathbf{P}(Y = y \,|X = x\,) \neq p\,(Y = y)$
   $\exists y, z : \mathbf{P}(Z = z \,|\,Y = y\,) \neq p\,(Z = z)$
2. Z and X: are (probably) dependent
3. $Z \perp X|Y$: are conditionally independent: $\forall x, y, z$

$$\mathbf{P}(Z = z \,|X = x, Y = y\,) = \mathbf{P}(Z = z \,|\,Y = y\,)$$

(if Y is fixed, then X and Z are independent)

Causality graph elements

$$X \leftarrow Y \rightarrow Z - \textbf{fork}$$

Example:

- X — ice cream sells
- Y — average temperature
- Z — crime ratio

Properties:

1. X and Y, Y and Z are dependent
2. X and Z are (probably) dependent
3. $X \perp Z | Y$ are conditionally independent

Causality graph elements

$$Y \to X \leftarrow Z \text{ — } \textbf{collider}$$

Example (illnes):

- X — bad symptoms
- Y — age
- Z — chronical diseases

Properties:

1. Y and X, Z and X are dependent
2. Y and Z are independent
3. $Y \not\perp Z | X$ are conditionally dependent

d-separation

The path P **is blocked** by Z, if:

1. P contains $A \to B \to C$, $A \leftarrow B \to C$, $B \in Z$
2. P contains $A \to B \leftarrow C$, $B \notin Z$ and all children of $B \notin Z$

If Z blocks all the paths from X to Y, then X and Y are **d-separated**:

$$X \perp Y | Z.$$

d-separation

The path P is **blocked** by Z, if:

1. P contains $A \to B \to C$, $A \leftarrow B \to C$, $B \in Z$
2. P contains $A \to B \leftarrow C$, $B \notin Z$ and all children of $B \notin Z$

If Z blocks all the paths from X to Y, then X and Y are **d-separated**.

Example:

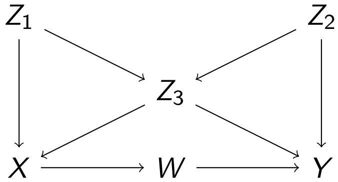

Pair	d-separation set
(Z_1, W)	X

d-separation

The path P is blocked by Z, if:

1. P contains $A \to B \to C$, $A \leftarrow B \to C$, $B \in Z$
2. P contains $A \to B \leftarrow C$, $B \notin Z$ and all children of $B \notin Z$

If Z blocks all the paths from X to Y, then X and Y are **d-separated**.

Example:

$$
\begin{tikzcd}
Z_1 \arrow[dd] \arrow[dr] & & Z_2 \arrow[dl] \arrow[dd] \\
& Z_3 \arrow[dl] \arrow[dr] & \\
X \arrow[r] & W \arrow[r] & Y
\end{tikzcd}
$$

Pair	d-separation set
(Z_1, W)	X
(Z_1, Y)	$\{Z_3, X, Z_2\}, \{Z_3, W, Z_2\}$

Model selection for Bayesian networks

- Generally, NP-hard problem
- Reduces to optimization problem with predefined search space or sampling problem
- Independence determination:
  - ML and MAP
  - Evidence
  - Information criteria

Markov random fields

Models are represented as undirected graphs.

Difference from Bayesian networks:

- No direction $\rightarrow$ cannot infer causality.
- The likelihood is factorized as follows:

$$p(\mathbf{x}) = \frac{1}{Z} \prod_C \psi(\mathbf{X}_C),$$

  where $\mathbf{X}_C$ is a maximal clicque, $\psi \geq 0$ is a potential function.
- Conditional indepdence: if all the paths from A to B go throught C, then $A \perp B | C$.

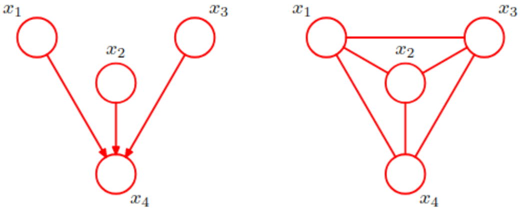

(Bishop)

Example: CRF and HMM

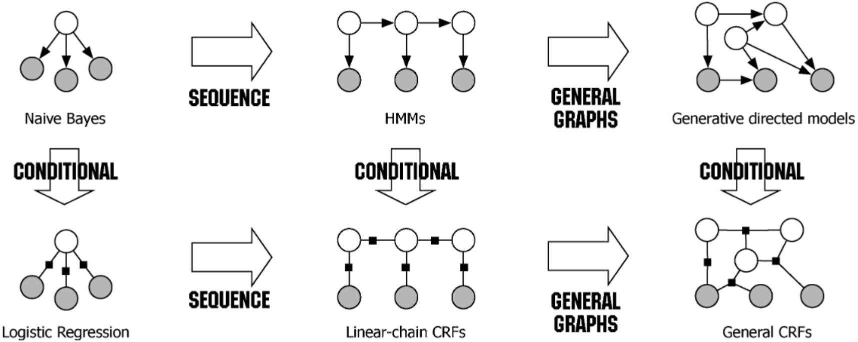

(Sutton, McCallum)

Inference in chains

(Bishop)

Naive likelihood calculation for x_n:

$$p(x_n) = \sum_{x_1} \sum_{x_2} \ldots, \sum_{x_{n-1}} \sum_{x_{n+1}} \cdots \sum_{x_N} p(\mathbf{x}),$$

For N discrete variables with K values the complexity is $O(K^N)$

Inference in chains: regroupping

$$p(x_n) = \sum_{x_1} \sum_{x_2} \dots, \sum_{x_{n-1}} \sum_{x_{n+1}} \cdots \sum_{x_N} p(\mathbf{x}),$$

$$p(\mathbf{x}) = \psi(x_1, x_2)\psi(x_2, x_3) \dots \psi(x_{N-1}, x_N).$$

Regroup the sum:

$$p(x_n) = \sum_{x_{n-1}} \psi(x_{n-1}, x_n) \dots \big(\sum_{x_1} \psi(x_1, x_2)\big) \times$$

$$\times \big(\sum_{x_{n+1}} \psi(x_n, x_{n+1}) \dots \big(\sum_{x_N} \psi(x_{N-1}, x_N)\big)\big).$$

Now complexity is $O(NK^2)$.

Message passing

$$p(x_n) = \underbrace{\sum_{x_{n-1}} \psi(x_{n-1}, x_n) \dots \left(\sum_{x_1} \psi(x_1, x_2)\right)}_{\mu_a(x_n)} \times$$

$$\times \underbrace{\left(\sum_{x_{n+1}} \psi(x_n, x_{n+1}) \dots \left(\sum_{x_N} \psi(x_{N-1}, x_N)\right)\right)}_{\mu_b(x_n)}.$$

Interpretation: $\mu_a(x_n)$ is a message transferred from x_{n-1} to x_n, $\mu_b(x_n)$ is a backward message from x_{n+1}.

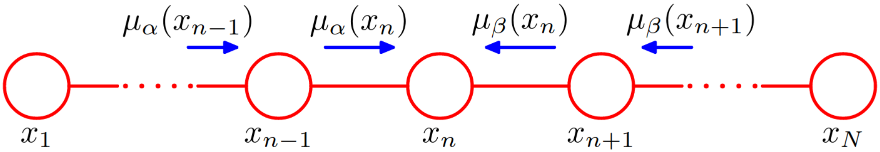

Inference in chains: details

The inference is iterative:

- calculate $\sum_{x_1} \psi(x_1, x_2) = \boldsymbol{\mu}_a(\mathbf{x}_2)$, that stores $\mu_a(x_2)$ for each value of x_2;
- calculate $\sum_{x_2} \psi(x_2, x_3)\big(\sum_{x_1} \psi(x_1, x_2)\big) = \sum_{x_2} \psi(x_2, x_3)\mu_a(x_2) = \boldsymbol{\mu}_a(\mathbf{x}_3)$;
- ...
- calculate$\sum_{x_{n+1}} \psi(x_n, x_{n+1})\mu_b(x_{n+1}) = \boldsymbol{\mu}_b(\mathbf{x}_n)$.
- for directed variables, where

$$\psi(x_1, x_2) = p(x_1)p(x_2|x_1), \quad \psi(x_i, x_{i+1}) = p(x_{i+1}|x_i),$$

μ_b should not be calculated:

$$\mu_b(x_n) = \sum_{x_{n+1}} \psi(x_n, x_{n+1}) \dots \big(\sum_{x_N} \psi(x_{N-1}, x_N)\big) =$$

$$= \sum_{x_{n+1}} p(x_{n+1}|x_n) \dots \big(\sum_{x_N} p(x_N|x_{N-1})\big) = 1.$$

Factor graph

Definition

Factor-graph is a bipartite graph with two types of vertivces: variables and factors.
The likelihood is a production of factors:

$$p(\mathbf{x}) = \prod_i f_i.$$

Example: model $p(x_1)p(x_2)p(x_3|x_2,x_1)$ has two variants of factorization:

$$f = p(x_1)p(x_2)p(x_3|x_2,x_1), \quad f_a = p(x_1), f_b = p(x_2), f_3 = p(x_1)p(x_2)p(x_3|x_2,x_1).$$

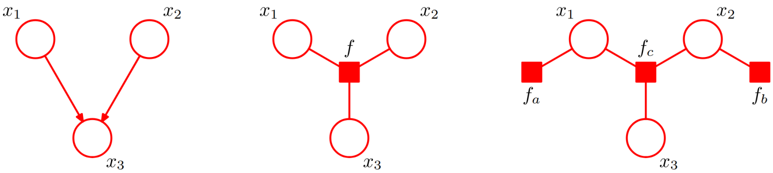

(Bishop)

Inference in factor-graphs: example

Sum-product: likelihood is a composition of messages from factors to variables.

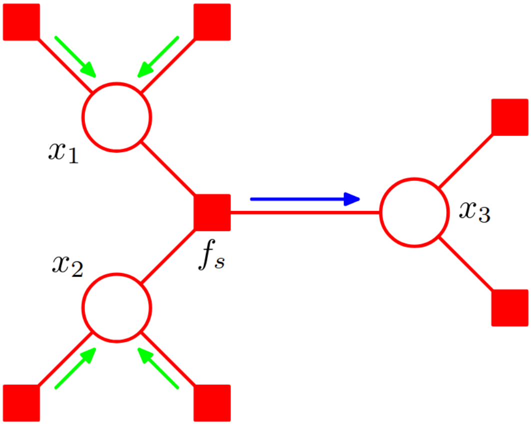

Model examples: RBM

$$p(\mathbf{x}, \mathbf{h}) = \frac{1}{Z}\exp(-E(\mathbf{x}, \mathbf{h})),$$

$$E = -\mathbf{w}_1^T\mathbf{x} - \mathbf{w}_2^T\mathbf{h} - \mathbf{x}^T\mathbf{W}_3\mathbf{h},$$

$p(\mathbf{h} = 1|\mathbf{x})$ and $p(\mathbf{x} = 1|\mathbf{h})$ are 1-layers with sigmoid activation.

Derivative of the log-likelihood for a single example x with respect to $\mathbf{w}$:

$$\frac{\partial \log p_{\mathbf{w}}(\mathbf{x})}{\partial \mathbf{w}} = \mathbb{E}_{\mathbf{x}',\mathbf{h}\sim p_{\mathbf{w}}(\mathbf{x})}\left[\frac{\partial E_{\mathbf{w}}(x')}{\partial \mathbf{w}}\right] - \frac{\partial E_{\mathbf{w}}(x)}{\partial \mathbf{w}}.$$

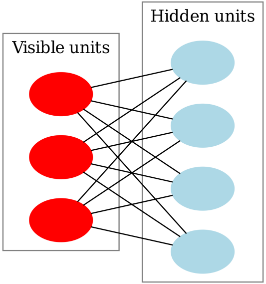

Model examples: Structured VAEs

Based on SLDS:

$$z_{t+1} | z_t \sim \pi^{t+1},$$

$$\mathbf{y}_t \sim \mathcal{N}(\mathbf{MLP}^{z_t}(\mathbf{x}_t)).$$

Optimization: optimize ELBO.
Inference: message-passing.

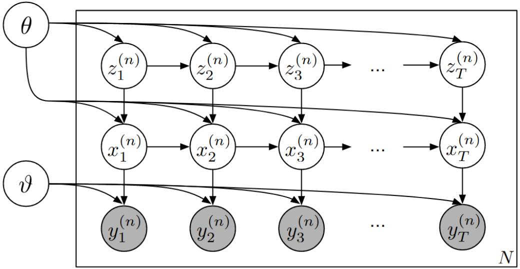

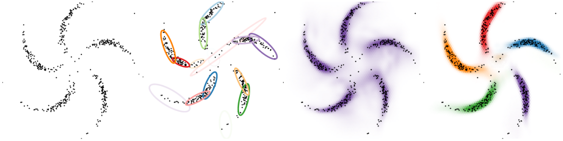

(a) Data (b) GMM (c) Density net (VAE) (d) GMM SVAE

References

- Bishop C. M. Pattern recognition //Machine learning. – 2006. – T. 128. – №. 9.
- Edwards D. Introduction to graphical modelling. – Springer Science & Business Media, 2012.
- Pearl J., Glymour M., Jewell N. P. Causal inference in statistics: A primer. – John Wiley & Sons, 2016.
- Hinton G. E., Salakhutdinov R. R. Reducing the dimensionality of data with neural networks //science. – 2006. – T. 313. – №. 5786. – C. 504-507.
- https://en.wikipedia.org/wiki/Restricted_Boltzmann_machine
- Johnson M. J. et al. Structured VAEs: Composing probabilistic graphical models and variational autoencoders //arXiv preprint arXiv:1603.06277. – 2016. – T. 2. – C. 2016.
- Johnson M. J. et al. Composing graphical models with neural networks for structured representations and fast inference //Advances in neural information processing systems. – 2016. – T. 29. – C. 2946-2954.
- Højsgaard, S., Edwards, D., Lauritzen, S. (2012). Graphical models with R. Springer Science & Business Media.
- Sutton, C, McCallum, A 2012, 'An Introduction to Conditional Random Fields', Foundations and Trends in Machine Learning, vol. 4, no. 4, pp. 267-373
- https://www.stat.umn.edu/geyer/5421/notes/graph.pdf

Organizational issues

- General requirements for library and algorithms:
  - Codestyle: at least PEP-8
  - The code must be commented
  - There must be an installation file (setup.py or alternative)
  - There must be requirements file with defined versions
  - All these points (and general work of library) will be checked offline
- Demo: must be available at the meeting (perfect case: ipynb or collab)
- Tests:
  - Coverage 75%;
  - Tests during build: not required
  - Will be checked offline
- Documentation: must be ready and deployed
- Blog-post: pre-final version must be ready
- Documentation and blog-post will be checked offline. For blog-post prepare for the external review.